

```
In [1]: from sklearn.preprocessing import LabelBinarizer
from sklearn.metrics import classification_report
from tensorflow.keras.models import Sequential
from tensorflow.keras.layers import Dense
from tensorflow.keras.optimizers import SGD
from tensorflow.keras.datasets import mnist
from tensorflow.keras import backend as K
import matplotlib.pyplot as plt
import numpy as np
import argparse
```

```
2023-09-13 11:11:58.759399: I tensorflow/tsl/cuda/cudart_stub.cc:28] Could not find cuda drivers on your machine, GPU will not be used.
2023-09-13 11:11:58.787163: I tensorflow/tsl/cuda/cudart_stub.cc:28] Could not find cuda drivers on your machine, GPU will not be used.
2023-09-13 11:11:58.787604: I tensorflow/core/platform/cpu_feature_guard.cc:182] This TensorFlow binary is optimized to use available CPU instructions in performance-critical operations.
To enable the following instructions: AVX2 FMA, in other operations, rebuild TensorFlow with the appropriate compiler flags.
2023-09-13 11:11:59.269270: W tensorflow/compiler/tf2tensorrt/utils/py_utils.cc:38] TF-TRT Warning: Could not find TensorRT
```

```
In [2]: print("[INFO] accessing MNIST...")
((trainX, trainY), (testX, testY)) = mnist.load_data()

[INFO] accessing MNIST...
```

```
In [3]: # each image in the MNIST dataset is represented as a 28x28x1
# image, but in order to apply a standard neural network we must
# first "flatten" the image to be simple list of 28x28=784 pixels
trainX = trainX.reshape((trainX.shape[0], 28 * 28 * 1))
testX = testX.reshape((testX.shape[0], 28 * 28 * 1))
```

```
In [4]: # scale data to the range of [0, 1]
trainX = trainX.astype("float32") / 255.0
testX = testX.astype("float32") / 255.0
```

```
In [5]: # convert the labels from integers to vectors
lb = LabelBinarizer()
trainY = lb.fit_transform(trainY)
testY = lb.transform(testY)
```

```
In [6]: # define the 784-256-128-10 architecture using Keras
model = Sequential()
model.add(Dense(256, input_shape=(784,), activation="sigmoid"))
model.add(Dense(128, activation="sigmoid"))
model.add(Dense(10, activation="softmax"))
```

```
In [7]: # train the model using SGD
print("[INFO] training network...")
sgd = SGD(0.01)
model.compile(loss="categorical_crossentropy", optimizer=sgd,
metrics=["accuracy"])
H = model.fit(trainX, trainY, validation_data=(testX, testY),
epochs=100, batch_size=128)
```

[INFO] training network...

Epoch 1/100

2023-09-13 11:12:01.320784: W tensorflow/tsl/framework/cpu_allocator_impl.cc:83] Allocation of 188160000 exceeds 10% of free system memory.

469/469 [=====] - 1s 2ms/step - loss: 2.27

12 - accuracy: 0.2246 - val_loss: 2.2328 - val_accuracy: 0.4521

Epoch 2/100

469/469 [=====] - 1s 2ms/step - loss: 2.19

82 - accuracy: 0.4281 - val_loss: 2.1520 - val_accuracy: 0.5051

Epoch 3/100

469/469 [=====] - 1s 2ms/step - loss: 2.10

46 - accuracy: 0.5502 - val_loss: 2.0372 - val_accuracy: 0.6329

Epoch 4/100

469/469 [=====] - 1s 2ms/step - loss: 1.96

82 - accuracy: 0.6148 - val_loss: 1.8712 - val_accuracy: 0.6407

Epoch 5/100

469/469 [=====] - 1s 2ms/step - loss: 1.78


```
In [8]: print("[INFO] evaluating network...")
predictions = model.predict(testX, batch_size=128)
print(classification_report(testY.argmax(axis=1),
predictions.argmax(axis=1),
target_names=[str(x) for x in lb.classes_]))
```

[INFO] evaluating network...

79/79 [=====] - 0s 1ms/step

	precision	recall	f1-score	support
0	0.94	0.98	0.96	980
1	0.97	0.97	0.97	1135
2	0.93	0.91	0.92	1032
3	0.91	0.92	0.91	1010
4	0.92	0.93	0.93	982
5	0.90	0.87	0.89	892
6	0.93	0.95	0.94	958
7	0.94	0.92	0.93	1028
8	0.90	0.89	0.90	974
9	0.91	0.91	0.91	1009
accuracy			0.93	10000
macro avg	0.92	0.92	0.92	10000
weighted avg	0.93	0.93	0.93	10000

```
In [9]: plt.style.use("ggplot")
plt.figure()
plt.plot(np.arange(0, 100), H.history["loss"], label="train_loss")
plt.plot(np.arange(0, 100), H.history["val_loss"], label="val_loss")
plt.plot(np.arange(0, 100), H.history["accuracy"], label="train_acc")
plt.plot(np.arange(0, 100), H.history["val_accuracy"], label="val_acc")
plt.title("Training Loss and Accuracy")
plt.xlabel("Epoch #")
plt.ylabel("Loss/Accuracy")
plt.legend()
plt.savefig("output.jpg")
```



```
In [10]: print("[INFO] accessing MNIST...")
((trainX, trainY), (testX, testY)) = mnist.load_data()

[INFO] accessing MNIST...
```

```
In [11]: trainX = trainX.reshape((trainX.shape[0], 28 * 28 * 1))
testX = testX.reshape((testX.shape[0], 28 * 28 * 1))
xtest_nonreshape=testX.reshape((testX.shape[0],28,28))
```

```
In [12]: # scale data to the range of [0, 1]
trainX = trainX.astype("float32") / 255.0
testX = testX.astype("float32") / 255.0
```

```
In [13]: # convert the labels from integers to vectors
lb = LabelBinarizer()
trainY = lb.fit_transform(trainY)
testY = lb.transform(testY)
```

```
In [14]: print("[INFO] training network...")
sgd = SGD(0.01)
model.compile(loss="categorical_crossentropy", optimizer=sgd,
metrics=["accuracy"])
H = model.fit(trainX, trainY, validation_data=(testX, testY),
epochs=100, batch_size=128)
```

```
[INFO] training network...
Epoch 1/100
```

```
2023-09-13 11:15:35.578222: W tensorflow/tsl/framework/cpu_allocator_impl.cc:83] Allocation of 188160000 exceeds 10% of free system memory.
```

```
469/469 [=====] - 1s 2ms/step - loss: 0.2654 - accuracy: 0.9235 - val_loss: 0.2601 - val_accuracy: 0.9254
```

```
Epoch 2/100
```

```
469/469 [=====] - 1s 2ms/step - loss: 0.2645 - accuracy: 0.9234 - val_loss: 0.2592 - val_accuracy: 0.9260
```

```
Epoch 3/100
```

```
469/469 [=====] - 1s 2ms/step - loss: 0.2637 - accuracy: 0.9237 - val_loss: 0.2586 - val_accuracy: 0.9261
```

```
Epoch 4/100
```

```
469/469 [=====] - 1s 2ms/step - loss: 0.2629 - accuracy: 0.9241 - val_loss: 0.2579 - val_accuracy: 0.9265
```

```
Epoch 5/100
```

```
469/469 [=====] - 1s 2ms/step - loss: 0.26
```

```
In [15]: print("[INFO] evaluating network...")
predictions = model.predict(testX, batch_size=128)
print(classification_report(testY.argmax(axis=1),
predictions.argmax(axis=1),
target_names=[str(x) for x in lb.classes_]))
```


```
[INFO] evaluating network...
```

```
79/79 [=====] - 0s 1ms/step
              precision    recall  f1-score   support
```

0	0.95	0.98	0.97	980
1	0.97	0.98	0.98	1135
2	0.94	0.94	0.94	1032
3	0.93	0.94	0.94	1010
4	0.93	0.95	0.94	982
5	0.93	0.90	0.92	892
6	0.94	0.96	0.95	958
7	0.95	0.93	0.94	1028
8	0.93	0.91	0.92	974
9	0.93	0.92	0.93	1009

accuracy			0.94	10000
macro avg	0.94	0.94	0.94	10000
weighted avg	0.94	0.94	0.94	10000

```
In [16]: plt.style.use("ggplot")
plt.figure()
plt.plot(np.arange(0, 100), H.history["loss"], label="train_loss")
plt.plot(np.arange(0, 100), H.history["val_loss"], label="val_loss")
plt.plot(np.arange(0, 100), H.history["accuracy"], label="train_acc")
plt.plot(np.arange(0, 100), H.history["val_accuracy"], label="val_acc")
plt.title("Training Loss and Accuracy")
plt.xlabel("Epoch #")
plt.ylabel("Loss/Accuracy")
plt.legend()
plt.savefig("output1.jpg")
```



```
In [17]: import random

for i in range(0,9):
    r=random.randint(0,100)
    plt.imshow(xtest_nonreshape[r])

prediction=model.predict(testX)
print(f"The predicted value is {np.argmax(prediction[r])}")
```

313/313 [=====] - 0s 585us/step
The predicted value is 7

