

```
In [8]: from sklearn.preprocessing import LabelBinarizer
from sklearn.metrics import classification_report
from tensorflow.keras.models import Sequential
from tensorflow.keras.layers import Dense
from tensorflow.keras.optimizers import SGD
from tensorflow.keras.datasets import cifar10
from tensorflow.keras import backend as K
import matplotlib.pyplot as plt
import numpy as np
import argparse
```

```
In [9]: print("[INFO] accessing MNIST...")
((trainX, trainY), (testX, testY)) = cifar10.load_data()

[INFO] accessing MNIST...
```

```
In [10]: # each image in the MNIST dataset is represented as a 28x28x1
# image, but in order to apply a standard neural network we must
# first "flatten" the image to be simple list of 28x28=784 pixels
trainX = trainX.reshape((trainX.shape[0], 32 * 32 * 3))
testX = testX.reshape((testX.shape[0], 32 * 32 * 3))
```

```
In [11]: # scale data to the range of [0, 1]
trainX = trainX.astype("float32") / 255.0
testX = testX.astype("float32") / 255.0
```

```
In [12]: # convert the labels from integers to vectors
lb = LabelBinarizer()
trainY = lb.fit_transform(trainY)
testY = lb.transform(testY)
```

```
In [15]: # define the 784-256-128-10 architecture using Keras
model = Sequential()
model.add(Dense(256, input_shape=(3072,), activation="sigmoid"))
model.add(Dense(128, activation="sigmoid"))
model.add(Dense(10, activation="softmax"))
```

```
In [16]: # train the model using SGD
print("[INFO] training network...")
sgd = SGD(0.01)
model.compile(loss="categorical_crossentropy", optimizer=sgd,
metrics=["accuracy"])
H = model.fit(trainX, trainY, validation_data=(testX, testY),
epochs=100, batch_size=256)
```

[INFO] training network...

Epoch 1/100

2023-09-13 12:19:30.991660: W tensorflow/tsl/framework/cpu_allocator_impl.cc:83] Allocation of 614400000 exceeds 10% of free system memory.

186/196 [=====>...] - ETA: 0s - loss: 2.3104 - accuracy: 0.1436

2023-09-13 12:19:41.969410: W tensorflow/tsl/framework/cpu_allocator_impl.cc:83] Allocation of 122880000 exceeds 10% of free system memory.

196/196 [=====] - 8s 8ms/step - loss: 2.3093 - accuracy: 0.1456 - val_loss: 2.2846 - val_accuracy: 0.1970
Epoch 2/100

196/196 [=====] - 1s 5ms/step - loss: 2.2774 - accuracy: 0.1957 - val_loss: 2.2690 - val_accuracy: 0.2263
Epoch 3/100

```
In [17]: print("[INFO] evaluating network...")
predictions = model.predict(testX, batch_size=128)
print(classification_report(testY.argmax(axis=1),
predictions.argmax(axis=1),
target_names=[str(x) for x in lb.classes_]))
```

[INFO] evaluating network...

57/79 [=====>.....] - ETA: 0s

2023-09-13 12:22:03.862180: W tensorflow/tsl/framework/cpu_allocator_impl.cc:83] Allocation of 122880000 exceeds 10% of free system memory.

79/79 [=====] - 0s 2ms/step

	precision	recall	f1-score	support
--	-----------	--------	----------	---------

0	0.43	0.49	0.45	1000
1	0.43	0.47	0.45	1000
2	0.28	0.22	0.25	1000
3	0.27	0.23	0.25	1000
4	0.38	0.27	0.31	1000
5	0.33	0.37	0.35	1000
6	0.39	0.44	0.41	1000
7	0.41	0.41	0.41	1000
8	0.45	0.53	0.49	1000
9	0.46	0.48	0.47	1000

accuracy			0.39	10000
macro avg	0.38	0.39	0.38	10000
weighted avg	0.38	0.39	0.38	10000

```
In [18]: plt.style.use("ggplot")
plt.figure()
plt.plot(np.arange(0, 100), H.history["loss"], label="train_loss")
plt.plot(np.arange(0, 100), H.history["val_loss"], label="val_loss")
plt.plot(np.arange(0, 100), H.history["accuracy"], label="train_acc")
plt.plot(np.arange(0, 100), H.history["val_accuracy"], label="val_acc")
plt.title("Training Loss and Accuracy")
plt.xlabel("Epoch #")
plt.ylabel("Loss/Accuracy")
plt.legend()
plt.savefig("output.jpg")
```

