

```
#import required libraries
import pandas as pd
import numpy as np
from scipy import stats
import tensorflow as tf
import matplotlib.pyplot as plt
import seaborn as sns
from pylab import rcParams
from sklearn.model_selection import train_test_split
from keras.models import Model, load_model
from keras.layers import Input, Dense
from keras.callbacks import ModelCheckpoint, TensorBoard
from keras import regularizers
from sklearn.preprocessing import StandardScaler
from sklearn.metrics import confusion_matrix, recall_score, accuracy_score, precision_score
```

```
sns.set(style='whitegrid', palette='muted', font_scale=1.5)
```

```
rcParams['figure.figsize'] = 14, 8
```

```
RANDOM_SEED = 42
```

```
LABELS = ["Normal", "Fraud"]
```

```
#upload the dataset
```

```
df = pd.read_csv("/content/creditcard.csv")
```

```
df.shape
```

```
(284807, 31)
```

```
df.isnull().values.any()
```

```
False
```

```
print(list(df.columns))
```

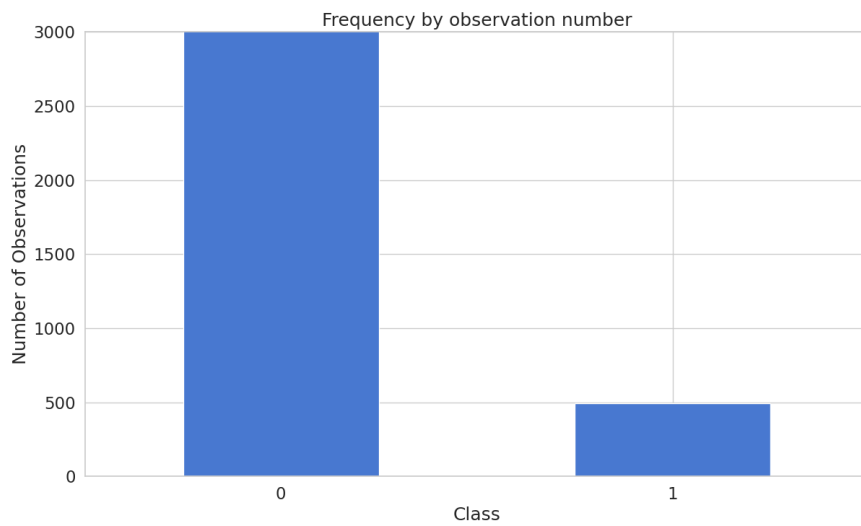
```
df.describe()
```

```
['Time', 'V1', 'V2', 'V3', 'V4', 'V5', 'V6', 'V7', 'V8', 'V9', 'V10', 'V11', 'V12', 'V13', 'V14', 'V15', 'V16', 'V17', 'V18', 'V19', 'V20', 'V21', 'V22', 'V23', 'V24', 'V25', 'V26', 'V27', 'V28', 'V29', 'V30', 'V31', 'Class']
```


	Time	V1	V2	V3	V4	V5
count	284807.000000	2.848070e+05	2.848070e+05	2.848070e+05	2.848070e+05	2.848070e+05
mean	94813.859575	1.168375e-15	3.416908e-16	-1.379537e-15	2.074095e-15	9.604095e-16
std	47488.145955	1.958696e+00	1.651309e+00	1.516255e+00	1.415869e+00	1.380211e+00
min	0.000000	-5.640751e+01	-7.271573e+01	-4.832559e+01	-5.683171e+00	-1.137418e+01
25%	54201.500000	-9.203734e-01	-5.985499e-01	-8.903648e-01	-8.486401e-01	-6.915625e-01
50%	84692.000000	1.810880e-02	6.548556e-02	1.798463e-01	-1.984653e-02	-5.433511e-01
75%	139320.500000	1.315642e+00	8.037239e-01	1.027196e+00	7.433413e-01	6.119200e-01
max	172792.000000	2.454930e+00	2.205773e+01	9.382558e+00	1.687534e+01	3.480134e+01

```
8 rows × 31 columns
```

```
#Visualizing the imbalanced dataset
count_classes = pd.value_counts(df['Class'], sort = True)
count_classes.plot(kind = 'bar', rot=0)
plt.xticks(range(len(df['Class'].unique())), df.Class.unique())
plt.ylim(0,3000)
plt.title("Frequency by observation number")
plt.xlabel("Class")
plt.ylabel("Number of Observations");
```



```
#Visualizing the imbalanced dataset
sub=df.iloc[:20000]
count_classes = pd.value_counts(sub['Class'], sort = True)
count_classes.plot(kind = 'pie', rot=0)
plt.xticks(range(len(sub['Class'].unique())), sub.Class.unique())
plt.legend()
plt.title("Frequency by observation number")
plt.xlabel("Class")
plt.ylabel("Number of Observations");
```



```
sc=StandardScaler()
df['Time'] = sc.fit_transform(df['Time'].values.reshape(-1, 1))
df['Amount'] = sc.fit_transform(df['Amount'].values.reshape(-1, 1))
```

```

raw_data = df.values
# The last element contains if the transaction is normal which is represented by a 0 and if fraud then 1
labels = raw_data[:, -1]
# The other data points are the electrocardiogram data
data = raw_data[:, 0:-1]
train_data, test_data, train_labels, test_labels = train_test_split(data, labels, test_size=0.2, random_state=2021)

```

```

#Normalize the data to have a value between 0 and 1
min_val = tf.reduce_min(train_data)
max_val = tf.reduce_max(train_data)
train_data = (train_data - min_val) / (max_val - min_val)
test_data = (test_data - min_val) / (max_val - min_val)
train_data = tf.cast(train_data, tf.float32)
test_data = tf.cast(test_data, tf.float32)

```

```

train_labels = train_labels.astype(bool)
test_labels = test_labels.astype(bool)
normal_train_data = train_data[~train_labels]
normal_test_data = test_data[~test_labels]
fraud_train_data = train_data[train_labels]
fraud_test_data = test_data[test_labels]
print(" No. of records in Fraud Train Data=",len(fraud_train_data))
print(" No. of records in Normal Train data=",len(normal_train_data))
print(" No. of records in Fraud Test Data=",len(fraud_test_data))
print(" No. of records in Normal Test data=",len(normal_test_data))

```

```

    No. of records in Fraud Train Data= 389
    No. of records in Normal Train data= 227456
    No. of records in Fraud Test Data= 103
    No. of records in Normal Test data= 56859

```

```

raw_data = df.values
# The last element contains if the transaction is normal which is represented by a 0 and if fraud then 1
labels = raw_data[:, -1]
# The other data points are the electrocardiogram data
data = raw_data[:, 0:-1]
train_data, test_data, train_labels, test_labels = train_test_split(data, labels, test_size=0.2, random_state=2021)

```

```

# Normalize the data to have a value between 0 and 1
min_val = tf.reduce_min(train_data)
max_val = tf.reduce_max(train_data)
train_data = (train_data - min_val) / (max_val - min_val)
test_data = (test_data - min_val) / (max_val - min_val)
train_data = tf.cast(train_data, tf.float32)
test_data = tf.cast(test_data, tf.float32)

```

```

train_labels = train_labels.astype(bool)
test_labels = test_labels.astype(bool)
normal_train_data = train_data[~train_labels]
normal_test_data = test_data[~test_labels]
fraud_train_data = train_data[train_labels]
fraud_test_data = test_data[test_labels]
print(" No. of records in Fraud Train Data=",len(fraud_train_data))
print(" No. of records in Normal Train data=",len(normal_train_data))
print(" No. of records in Fraud Test Data=",len(fraud_test_data))
print(" No. of records in Normal Test data=",len(normal_test_data))

```

```

    No. of records in Fraud Train Data= 389
    No. of records in Normal Train data= 227456
    No. of records in Fraud Test Data= 103
    No. of records in Normal Test data= 56859

```

```

nb_epoch = 50
batch_size = 64
input_dim = normal_train_data.shape[1] #num of columns, 30
encoding_dim = 14
hidden_dim_1 = int(encoding_dim / 2) #
hidden_dim_2=4
learning_rate = 1e-7

```

```

#Encoder converts it into latent representation
#input Layer
input_layer = tf.keras.layers.Input(shape=(input_dim, ))
#Encoder
encoder = tf.keras.layers.Dense(encoding_dim, activation="tanh",
activity_regularizer=tf.keras.regularizers.l2(learning_rate))(input_layer)
encoder=tf.keras.layers.Dropout(0.2)(encoder)

```

```
encoder = tf.keras.layers.Dense(hidden_dim_1, activation='relu')(encoder)
encoder = tf.keras.layers.Dense(hidden_dim_2, activation=tf.nn.leaky_relu)(encoder)
```

```
#Decoder networks convert it back to the original
```

```
# Decoder
```

```
decoder = tf.keras.layers.Dense(hidden_dim_1, activation='relu')(encoder)
decoder=tf.keras.layers.Dropout(0.2)(decoder)
decoder = tf.keras.layers.Dense(encoding_dim, activation='relu')(decoder)
decoder = tf.keras.layers.Dense(input_dim, activation='tanh')(decoder)
```

```
#Autoencoder
```

```
autoencoder = tf.keras.Model(inputs=input_layer, outputs=decoder)
autoencoder.summary()
```

```
Model: "model"
```

Layer (type)	Output Shape	Param #
input_1 (InputLayer)	[(None, 30)]	0
dense (Dense)	(None, 14)	434
dropout (Dropout)	(None, 14)	0
dense_1 (Dense)	(None, 7)	105
dense_2 (Dense)	(None, 4)	32
dense_3 (Dense)	(None, 7)	35
dropout_1 (Dropout)	(None, 7)	0
dense_4 (Dense)	(None, 14)	112
dense_5 (Dense)	(None, 30)	450

```
=====  
Total params: 1168 (4.56 KB)  
Trainable params: 1168 (4.56 KB)  
Non-trainable params: 0 (0.00 Byte)  
=====
```

```
cp = tf.keras.callbacks.ModelCheckpoint(filepath="autoencoder_fraud.h5",mode='min', monitor='val_loss', verbose=2, save_best_only=True)
```

```
# define our early stopping
```

```
early_stop = tf.keras.callbacks.EarlyStopping(  
    monitor='val_loss',  
    min_delta=0.0001,  
    patience=10,  
    verbose=1,  
    mode='min',  
    restore_best_weights=True)
```

```
#Compile the Autoencoder
```

```
autoencoder.compile(metrics=['accuracy'],  
loss='mean_squared_error',
```

```
optimizer='adam')
```

```
#Train the Autoencoder
```

```
history = autoencoder.fit(normal_train_data, normal_train_data,  
epochs=nb_epoch,  
batch_size=batch_size,  
shuffle=True,
```

```
validation_data=(test_data, test_data),  
verbose=1,
```

```
callbacks=[cp, early_stop]  
)
```

```
Epoch 1/50  
3535/3554 [=====>.] - ETA: 0s - loss: 0.0048 - accuracy: 0.0450  
Epoch 1: val_loss improved from inf to 0.00002, saving model to autoencoder_fraud.h5  
3554/3554 [=====] - 10s 2ms/step - loss: 0.0048 - accuracy: 0.0451 - val_loss: 1.9893e-05 - val_accuracy: 0.0451  
Epoch 2/50  
89/3554 [.....] - ETA: 6s - loss: 1.9877e-05 - accuracy: 0.0571/usr/local/lib/python3.10/dist-packages/keras/saving_api.save_model(  
3546/3554 [=====>.] - ETA: 0s - loss: 1.9335e-05 - accuracy: 0.0695  
Epoch 2: val_loss did not improve from 0.00002  
3554/3554 [=====] - 10s 3ms/step - loss: 1.9331e-05 - accuracy: 0.0696 - val_loss: 2.0112e-05 - val_accuracy: 0.0696  
Epoch 3/50  
3552/3554 [=====>.] - ETA: 0s - loss: 1.9447e-05 - accuracy: 0.0616
```

```

Epoch 3: val_loss did not improve from 0.00002
3554/3554 [=====] - 9s 3ms/step - loss: 1.9445e-05 - accuracy: 0.0616 - val_loss: 2.0113e-05 - val_accuracy
Epoch 4/50
3540/3554 [=====>.] - ETA: 0s - loss: 1.9542e-05 - accuracy: 0.0636
Epoch 4: val_loss did not improve from 0.00002
3554/3554 [=====] - 7s 2ms/step - loss: 1.9550e-05 - accuracy: 0.0636 - val_loss: 2.0032e-05 - val_accuracy
Epoch 5/50
3545/3554 [=====>.] - ETA: 0s - loss: 1.9540e-05 - accuracy: 0.0609
Epoch 5: val_loss did not improve from 0.00002
3554/3554 [=====] - 8s 2ms/step - loss: 1.9545e-05 - accuracy: 0.0609 - val_loss: 2.0122e-05 - val_accuracy
Epoch 6/50
3551/3554 [=====>.] - ETA: 0s - loss: 1.9547e-05 - accuracy: 0.0587
Epoch 6: val_loss did not improve from 0.00002
3554/3554 [=====] - 8s 2ms/step - loss: 1.9549e-05 - accuracy: 0.0587 - val_loss: 2.0195e-05 - val_accuracy
Epoch 7/50
3544/3554 [=====>.] - ETA: 0s - loss: 1.9520e-05 - accuracy: 0.0594
Epoch 7: val_loss did not improve from 0.00002
3554/3554 [=====] - 7s 2ms/step - loss: 1.9533e-05 - accuracy: 0.0594 - val_loss: 2.0377e-05 - val_accuracy
Epoch 8/50
3535/3554 [=====>.] - ETA: 0s - loss: 1.9518e-05 - accuracy: 0.0596
Epoch 8: val_loss did not improve from 0.00002
3554/3554 [=====] - 8s 2ms/step - loss: 1.9527e-05 - accuracy: 0.0600 - val_loss: 2.0062e-05 - val_accuracy
Epoch 9/50
3554/3554 [=====] - ETA: 0s - loss: 1.9528e-05 - accuracy: 0.0581
Epoch 9: val_loss did not improve from 0.00002
3554/3554 [=====] - 8s 2ms/step - loss: 1.9528e-05 - accuracy: 0.0581 - val_loss: 2.0419e-05 - val_accuracy
Epoch 10/50
3541/3554 [=====>.] - ETA: 0s - loss: 1.9514e-05 - accuracy: 0.0601
Epoch 10: val_loss did not improve from 0.00002
3554/3554 [=====] - 8s 2ms/step - loss: 1.9513e-05 - accuracy: 0.0600 - val_loss: 2.0367e-05 - val_accuracy
Epoch 11/50
3537/3554 [=====>.] - ETA: 0s - loss: 1.9505e-05 - accuracy: 0.0618
Epoch 11: val_loss did not improve from 0.00002
Restoring model weights from the end of the best epoch: 1.
3554/3554 [=====] - 8s 2ms/step - loss: 1.9508e-05 - accuracy: 0.0617 - val_loss: 2.0001e-05 - val_accuracy
Epoch 11: early stopping

```

```

#Plot training and test loss
plt.plot(history['loss'], linewidth=2, label='Train')
plt.plot(history['val_loss'], linewidth=2, label='Test')
plt.legend(loc='upper right')
plt.title('Model loss')
plt.ylabel('Loss')
plt.xlabel('Epoch')
#plt.ylim(ymin=0.70,ymax=1)
plt.show()

```


```

test_x_predictions = autoencoder.predict(test_data)
mse = np.mean(np.power(test_data - test_x_predictions, 2), axis=1)
error_df = pd.DataFrame({'Reconstruction_error': mse,
                          'True_class': test_labels})

1781/1781 [=====] - 2s 1ms/step

threshold_fixed = 0
groups = error_df.groupby('True_class')
fig, ax = plt.subplots()
for name, group in groups:
    ax.plot(group.index, group.Reconstruction_error, marker='o', ms=3.5, linestyle='', label= "Fraud" if name == 1 else "Normal")
ax.hlines(threshold_fixed, ax.get_xlim()[0], ax.get_xlim()[1], colors="r", zorder=100, label='Threshold')
ax.legend()
plt.title("Reconstruction error for normal and fraud data")
plt.ylabel("Reconstruction error")
plt.xlabel("Data point index")
plt.show()

```



```

threshold_fixed = 0.0001
y_pred = [1 if e > threshold_fixed else 0 for e in error_df.Reconstruction_error.values]
conf_matrix = confusion_matrix(error_df.True_class, y_pred)

plt.figure(figsize=(12, 12))
sns.heatmap(conf_matrix, xticklabels=LABELS, yticklabels=LABELS, annot=True, fmt="d");
plt.title("Confusion matrix")
plt.ylabel('True class')
plt.xlabel('Predicted class')
plt.show()

```

